

Post-Session Assessment & Feedback

Session: Build Your Own Clinical AI Agent

CME: Artificial Intelligence in Healthcare — From Theory to Clinical Mastery

SIMS&RH Tumkur · 12 August 2026 · 12:30–13:15

Faculty: Dr. Siddalingaiah H S, Professor, Dept. of Community Medicine

For KMC credit hour documentation.

Name: __ **Designation:** __ **Dept:** __

Part A — Assessment (5 questions, single best answer)

1. Which of the following best distinguishes an AI *agent* from an AI *chatbot*?

- A. An agent uses a larger language model
 - B. An agent pursues a goal over multiple steps and can use tools, deciding its own next step within limits set by the user
 - C. An agent is trained on medical data specifically
 - D. An agent does not make mistakes
-

2. A clinician builds a guideline assistant and asks it a question that the uploaded guidelines do not address. The tool produces a confident, fluent, plausible answer. This most likely indicates:

- A. The tool has additional medical knowledge and is being helpful
 - B. The tool lacks an effective grounding guardrail and is generating an ungrounded answer
 - C. The guidelines were uploaded incorrectly
 - D. The model needs to be retrained
-

3. Which statement about safety rules in a clinical AI agent is correct?

- A. Safety rules written into the prompt are enforced by the model and cannot be bypassed
 - B. Safety rules implemented in the surrounding application code are more reliably enforced than rules written into the prompt
 - C. Safety rules are unnecessary if the model is recent enough
 - D. Safety rules should be added only after the tool is deployed
-

4. Before pasting clinical text into a public AI service, the *minimum* requirement is:

- A. Informed consent from the patient for AI use
 - B. Removal of identifiers — name, contact details, hospital number, address, full date of birth
 - C. Approval from the hospital's IT department
 - D. Nothing, provided the clinician does not save the output
-

5. A vendor demonstrates an AI tool for clinical decision support. Which single request is most informative about its safety?

- A. "Show me your accuracy figures"
 - B. "Show me the list of hospitals using it"
 - C. "Show me it refusing to answer a question outside its knowledge"
 - D. "Show me which model it uses"
-

Part B — Feedback

1. Rate this session (1 = poor, 5 = excellent)

	1	2	3	4	5
Relevance to my work	☐	☐	☐	☐	☐
Clarity of explanation	☐	☐	☐	☐	☐
Usefulness of the live demonstration	☐	☐	☐	☐	☐
Practical takeaways I can use	☐	☐	☐	☐	☐
Pace and time management	☐	☐	☐	☐	☐

2. How likely are you to build an AI agent yourself in the next month?

☐ Already planning one ☐ Likely ☐ Unsure ☐ Unlikely

3. What is the one thing from this session you will actually use?

4. What did you want more of? What would you cut?

5. Would you attend a hands-on workshop where you build one yourself?

☐ Yes, half day ☐ Yes, full day ☐ No

6. Anything else:

Answer Key — For Faculty Use

Q	Answer	Teaching point
1	B	The defining feature is the <i>loop</i> : multi-step, tool-using, self-directed within limits. Model size and training data are irrelevant to the distinction. D is false of any AI system.
2	B	This is ungrounded generation — "hallucination". Fluency and confidence are not correlated with correctness, which is precisely why a refusal mechanism matters.
3	B	Prompt-based rules are requests the model may or may not honour. Rules enforced in surrounding code execute deterministically. Demonstrated live via the pre-model grounding threshold.
4	B	De-identification is the minimum. C may also be required institutionally, and consent (A) becomes relevant for wider deployment — but neither substitutes for removing identifiers first.
5	C	Accuracy figures are vendor-selected; deployment lists say nothing about safety. A demonstrated refusal is the single hardest thing to fake and the most informative about grounding.

Suggested pass mark: 4 of 5.

If a cohort scores poorly on Q3, it is worth re-emphasising in future sessions — it is the least intuitive point and the most consequential for anyone actually procuring a tool.