

The Clinical AI Safety Checklist

Ten questions to ask before any AI tool touches your patients — whether you built it or a vendor is selling it to you.

Print this. Take it to the meeting.

1 · Can it say "I don't know"?

Ask it something outside its knowledge. **Does it refuse, or does it guess?**

A tool that never refuses is guessing and not telling you. If a vendor cannot demonstrate a refusal on demand, that is your answer.

☐ Demonstrated ☐ Failed ☐ Not tested

2 · Where did that answer come from?

Every clinical claim should point to a source you can open and check.

"The AI said so" is not a source. If you cannot trace a statement back to a document, you cannot defend the decision built on it.

☐ Cites checkable sources ☐ No sources ☐ Sources given but wrong

3 · Who is clinically responsible for its output?

A **named** person, not a department or a committee.

If nobody's name is on it, nobody is checking it. This is the question that most often has no good answer.

Named owner: ____

4 · What happens to the data we put in?

Is it stored? For how long? Is it used to train the model? Which country is it in?

Under the **Digital Personal Data Protection Act 2023**, you have duties as a data fiduciary. "The vendor handles that" is not a defence.

☐ Answered in writing ☐ Vague ☐ Not answered

5 · Has identifiable patient data been excluded — and how do you know?

Not "we told staff not to." **What automatically stops it?**

Names, phone numbers, hospital numbers, addresses, full dates of birth, photographs. A check that depends on people remembering will fail on a busy day.

☐ Automated check ☐ Policy only ☐ Nothing

6 · How well does it actually perform — measured against what?

Accuracy on what sample, judged by whom, compared to which standard?

Vendor demos are chosen to succeed. Ask for performance on cases the vendor did not pick, ideally yours.

☐ Independent evaluation ☐ Vendor figures only ☐ No data

7 · How does it fail?

Not *whether* — *how*. Does it fail loudly or silently?

A tool that fails loudly is safe. A tool that quietly returns a plausible wrong answer is dangerous, and the more fluent it is, the more dangerous.

☐ Fails visibly ☐ Fails silently ☐ Unknown

8 · Can a human override it, and is that easy?

The override must be at least as easy as accepting.

If disagreeing takes four extra clicks and accepting takes none, the tool is making the decisions regardless of what the policy says.

☐ Easy override ☐ Difficult ☐ None

9 · Is there an audit trail?

Who asked what, when, what it answered, and what the clinician did with it.

You need this for incident review, for medico-legal defence, and for knowing whether the thing is being used at all. Build it in from day one — it cannot be added retrospectively.

☐ Full log ☐ Partial ☐ None

10 · What is it *not* for — written down, agreed, and communicated?

The boundary must exist on paper before the tool is in use.

Every AI incident in healthcare so far has involved a tool used slightly outside what it was validated for. The boundary you don't write down is the one that gets crossed.

Not to be used for: _____

Scoring

Answered well	What it means
9–10	Proceed, with monitoring
6–8	Pilot only. Fix the gaps before wider use.
3–5	Not ready. Do not put this in front of patients.
0–2	Walk away.

The three questions that matter most

If you only have five minutes in that meeting, ask these:

1. **"Show me it refusing to answer something."** (*Question 1*)
2. **"Whose name is on the output?"** (*Question 3*)
3. **"What is this NOT for?"** (*Question 10*)

The quality of the pause before each answer tells you most of what you need to know.

Prepared for the CME on Artificial Intelligence in Healthcare, SIMS&RH Tumkur, 12 August 2026 — session: Build Your Own Clinical AI Agent, Dr. Siddalingaiah H S, Dept. of Community Medicine.